

# Class 4

New York University

**Summer 2017**



1. Review
2. Project Discussion
3. Project Breakout
4. Spark RDD Design
5. Scala Tuples in Spark

1. Review
2. **Project Discussion** (*see [BDAD\\_SU17\\_ProjectInfo.pdf](#)*)
3. Project Breakout
4. Spark RDD Design
5. Scala Tuples in Spark

1. Review
2. Project Discussion
3. **Project Breakout**
4. Spark RDD Design
5. Scala Tuples in Spark

1. Review
2. Project Discussion
3. Project Breakout
4. **Spark RDD Design**
5. Scala Tuples in Spark

- Spark processing involves one of the following actions
  - Creating new RDDs
  - Transforming existing RDDs
  - Calling operations on RDDs to compute a result
- Under the hood, Spark distributes the data contained in RDDs
  - Data is distributed automatically to cluster servers
  - The operations you specify are parallelized across the cluster too

- An RDD is an immutable, distributed collection of objects
- An RDD contains
  - Python, Scala, or Java objects
  - It can even contain objects from user-defined classes

- Each RDD object is referred to as a *partition*
  - Specified computations are performed on each partition
  - Partitions can exist on any of the worker nodes in the cluster



- RDDs are created by
  - Loading an external dataset
  - Distributing a collection across the cluster
  - Transforming an existing RDD into a new RDD

- As discussed in an earlier class, RDDs offer two types of operations
  - *Transformations* construct a new RDD from an existing one
  - *Actions* compute a result based on processing an RDD
    - Actions do not create a new RDD
    - Actions return a result to the driver program or write the result to storage

- Transformations are performed lazily
  - I.e., when an action is issued, this is when all transformations are executed
  - This allows Spark to optimize operations to reduce the amount of data processed whenever possible
  - For example, `first()` causes Spark to read the first line only, not the entire (potentially Big Data) file
    - Contrast with MapReduce, where the entire file would be read

- Spark recomputes the related RDDs each time you run an action on them
  - You can avoid recomputing by using `RDD.persist()` (or `cache()`)
- After the first computation has been performed, the RDD can be persisted to memory (for example)
  - The RDD is persisted as partitions across the cluster
  - Once persisted, the data is available for future operations without having to be recomputed
- RDDs can also be persisted to disk

- Data is not persisted by default
  - Spark allows the programmer to control what should be persisted, and what need not be persisted
  - If data were persisted by default, there would be times when storage and time would be wasted to store data even though it would only be used once!
- We will look at options for persisting in a future class

- What is the number of partitions created?
  - By default, the number of partitions created for an RDD created from an HDFS file is the number of blocks in the file
  - If you want to verify the number of partitions created, use the `partitions` method of RDDs

- To check the number of partitions:

```
scala> someRDD.partitions.size
```

```
res0: Int = 30
```

- Here is an example that creates 30 partitions:

```
someRDD = sc.parallelize(range(101), 30)
```

1. Review
2. Project Discussion
3. Project Breakout
4. Spark RDD Design
5. **Scala Tuples in Spark**



## A Tuple is

- A group of individual values that can be treated as a single entity
- A compile-time entity

## Common uses of Tuples

- For returning more than one value from a function
- Key-value pairs
- For sending multiple values in a single message between concurrent processes
- For buffering a data record / related data of varying types

## Limitations of Tuples

- The number of values in a Tuple cannot be changed after it is initialized
- Tuples consist of between two (min) and twenty-two (max) values

A tuple containing two fields is a Tuple2, or *pair*

- Explicit declaration of a Tuple2 variable

```
val myTup2A = Tuple2(4, "iFruit")  
> myTup2A: (Int, String) = (4,iFruit)  
  
myTup2A.getClass  
> Class[_ <: (Int, String)] = class scala.Tuple2
```

- The -> syntax is available for Tuple2, but not for larger tuples

```
val myTup2B = 4 -> "iFruit"  
> myTup2B: (Int, String) = (4,iFruit)
```

- You can also allow the type to be inferred

```
val myTup2C = (4, "iFruit")  
> myTup2C: (Int, String) = (4,iFruit)
```

- Tuples contain two or more fields
- Individual tuple values can be accessed with `_1`, `_2` syntax
- Notice that tuples are *one*-based

```
val myTup2B = 4 -> "iFruit"  
> myTup2B: (Int, String) = (4,iFruit)  
  
myTup2B._1  
> Int = 4  
  
myTup2B._2  
> String = iFruit
```

- *swap* is syntactic sugar that works for Tuple2 only

```
myTup2B.swap  
> (String, Int) = (iFruit,4)
```

- Tuples with more than 2 elements: **TupleN**
- Declaring **TupleN** variables is similar to declaring **tuple2** variables
- Use **\_n** to access values in a **TupleN**

```
val myTup = (4,"MeToo","1.0",37.5,41.3,"Enabled")
> myTup: (Int, String, String, Double, Double, String) =
    (4,MeToo,1.0,37.5,41.3,Enabled)

myTup.getClass
> Class[_ <: (Int, String, String, Double, Double, String)] =
    class scala.Tuple6

println( myTup._3 + " / " + myTup._5 )
> 1.0 / 41.3
```



- Use **productPrefix** to get the tuple's class name as a string
- Use **productArity** to get the tuple size as an integer

```
val oneRecord = ("2014-03-15:10:10:20", "MeeToo", 3.0,  
"8316b507-7620-47aa-b56b-cae5cb2cd819", 0, 19, 69, 31, 51, 44, "TRUE",  
"enabled", "disabled", 33.4467594, -111.3653269)  
> oneRecord: (String, String, Double, String, Int, Int, Int, Int, Int,  
Int, String, String, String, Double, Double) =  
(2014-03-15:10:10:20,MeeToo,3.0,8316b507-7620-47aa-b56b-cae5cb2cd819,  
0,19,69,31,51,44,TRUE,enabled,disabled,33.4467594,-111.3653269)  
  
oneRecord.productPrefix  
> String = Tuple15  
  
oneRecord.productArity  
> Int = 15
```

- The values in a tuple can be converted to a single string using **toString**
  - Note that parentheses and commas are part of the new string

```
oneRecord._4  
> String = 8316b507-7620-47aa-b56b-cae5cb2cd819  
  
oneRecord.toString  
> String = (2014-03-15:10:10:20,MeeToo,3.0,8316b507-7620-47aa-b56b-  
cae5cb2cd819,  
0,19,69,31,51,44,TRUE,enabled,disabled,33.4467594,-111.3653269)
```

- Convert a string to a tuple using **partition**
- The **partition** function
  - Accepts a condition
    - In the example below, the condition is 'isUpper'
  - Returns a Tuple2 where the first value is what satisfied the condition
    - The second value contains what did not satisfy the condition

```
val myTup = ("Oranges", "Bananas", "apples", "Guavas")

val myStr = myTup.toString
> myStr: String =
  (Oranges,Bananas,apples,Guavas)

val myTup2 = myStr.partition(_.isUpper)
> (String, String) = (OBG, (ranges,ananas,apples,uavas))

val sortedValues = myTup2._1.sorted
> String = BGO
```

## Homework

See homework packet.