

Class 6: Pair RDDs

New York University

Summer 2017



- **Key-Value Pair RDDs**
- Map-Reduce
- Other Pair RDD Operations

■ What are Pair RDDs?

- Each RDD element must be a key-value pair (a two-element tuple)
- Keys and values can be any type
- Use with map-reduce algorithms
- Many additional functions are available for common data processing needs
 - e.g., sorting, joining, grouping, counting, etc.

Pair RDD

(key1, value1)
(key2, value2)
(key3, value3)
...

- **Common functions for creating Pair RDDs**
 - **map**
 - **flatMap / flatMapValues**
 - **keyBy**

- Example: Create a Pair RDD from a tab-separated file

```
> val users = sc.textFile(file) \
    .map(line => line.split('\t')) \
    .map(fields => (fields(0), fields(1)))
```

```
user001\tFred Flintstone
user090\tBugs Bunny
user111\tHarry Potter
...
```



(user001, Fred Flintstone)
(user090, Bugs Bunny)
(user111, Harry Potter)
...

```
> sc.textFile(logfile) \  
  .keyBy(line => line.split(' ')(2))
```

User ID

```
56.38.234.188 - 99788 "GET /KBDOC-00157.html HTTP/1.0" ...  
56.38.234.188 - 99788 "GET /theme.css HTTP/1.0" ...  
203.146.17.59 - 25254 "GET /KBDOC-00230.html HTTP/1.0" ...  
...
```

```
(99788,56.38.234.188 - 99788 "GET /KBDOC-00157.html...)  
(99788,56.38.234.188 - 99788 "GET /theme.css...)  
(25254,203.146.17.59 - 25254 "GET /KBDOC-00230.html...)  
...
```

■ How would you do this?

- Input: a list of postal codes with latitude and longitude
- Output: **postal code** (key) and **lat/long** pair (value)

00210	43.005895	-71.013202
00211	43.005895	-71.013202
00212	43.005895	-71.013202
00213	43.005895	-71.013202
00214	43.005895	-71.013202
...		



(00210, (43.005895, -71.013202))
(00211, (43.005895, -71.013202))
(00212, (43.005895, -71.013202))
(00213, (43.005895, -71.013202))
...

```
> sc.textFile("file").  
  map(line => line.split('\t')).  
  map(fields => (fields(0), fields(1), fields(2)))
```

```
00210 43.005895 -71.013202  
00211 43.005895 -71.013202  
00212 43.005895 -71.013202  
00213 43.005895 -71.013202  
00214 43.005895 -71.013202
```

...



```
(00210, (43.005895, -71.013202))  
(00211, (43.005895, -71.013202))  
(00212, (43.005895, -71.013202))  
(00213, (43.005895, -71.013202))  
...
```


- Key-Value Pair RDDs
- **Map-Reduce**
- Other Pair RDD Operations

- **Map-reduce is a common programming model**
 - Easily applicable to distributed processing of large data sets
- **Hadoop MapReduce is a well known implementation**
 - Somewhat limited
 - Each job has one Map phase, one Reduce phase
 - Job output is saved to files
- **Spark implements map-reduce with much greater flexibility**
 - Map and reduce functions can be interspersed
 - Results can be stored in memory
 - Operations can easily be chained

- **Map-reduce in Spark works on Pair RDDs**
- **Map phase**
 - Operates on one record at a time
 - Maps each record to one or more new records
 - e.g. `map`, `flatMap`, `filter`, `keyBy`
- **Reduce phase**
 - Works on map output, or output from other tools in same key-value format
 - Consolidates multiple records
 - e.g. `reduceByKey`, `sortByKey`, `mean`

■ How would you do this?

- Input: Text lines
- Output: Word count

Input Data

the cat sat on the mat
the aardvark sat on the sofa



Output

aardvark	1
cat	1
mat	1
on	2
sat	2
sofa	1
the	4

■ Use flatMap

Input Data

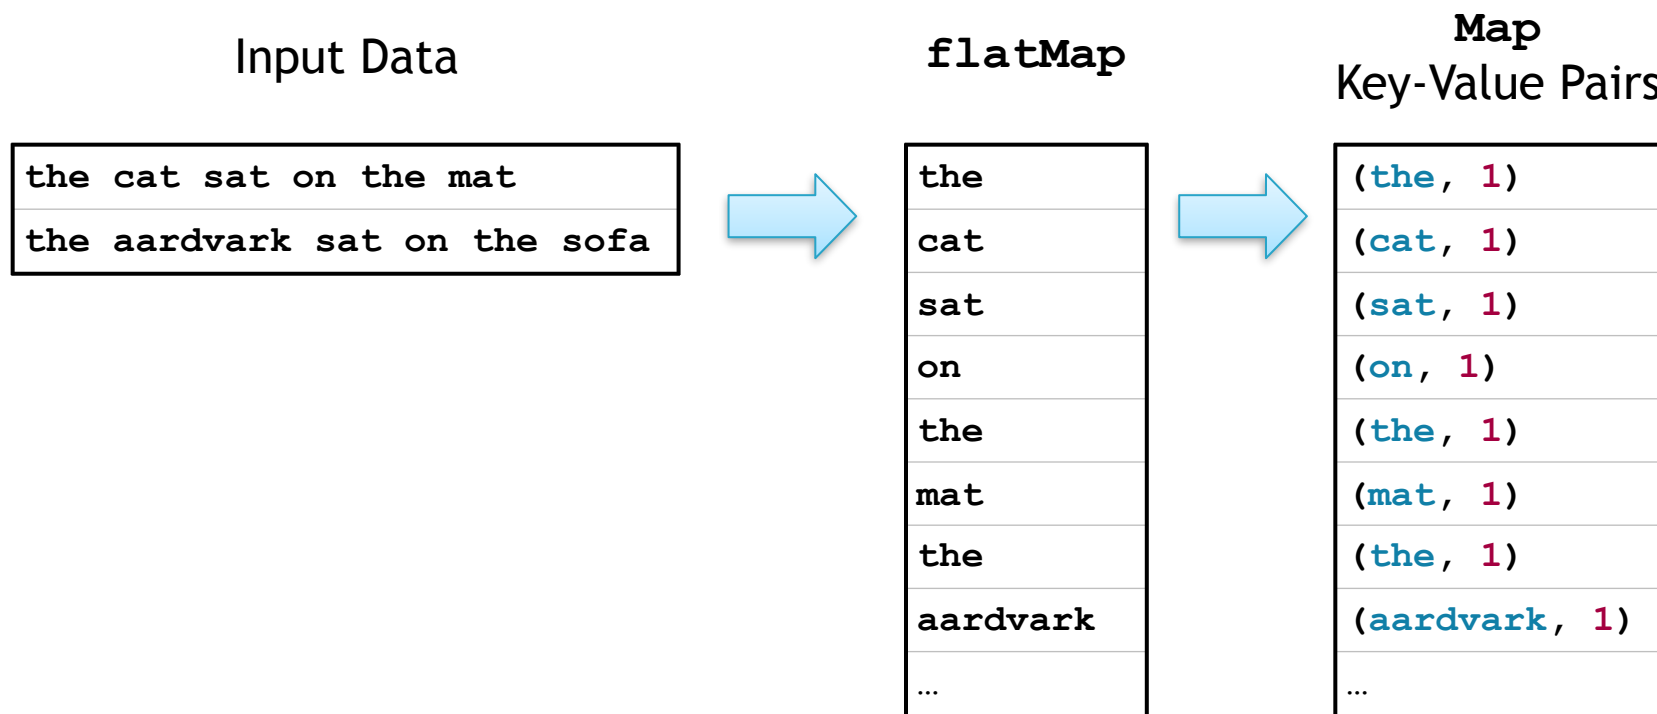
the cat sat on the mat
the aardvark sat on the sofa



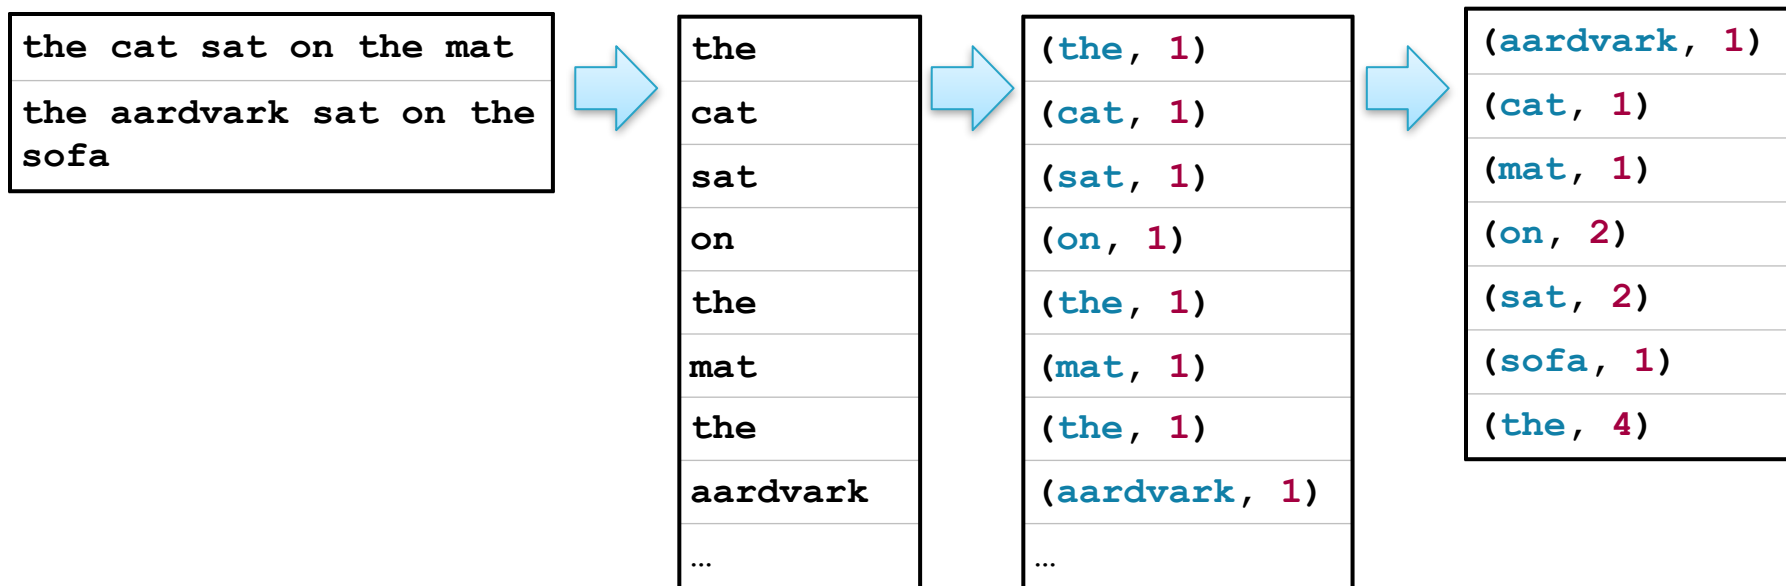
Output

the
cat
sat
on
the
mat
the
aardvark
...

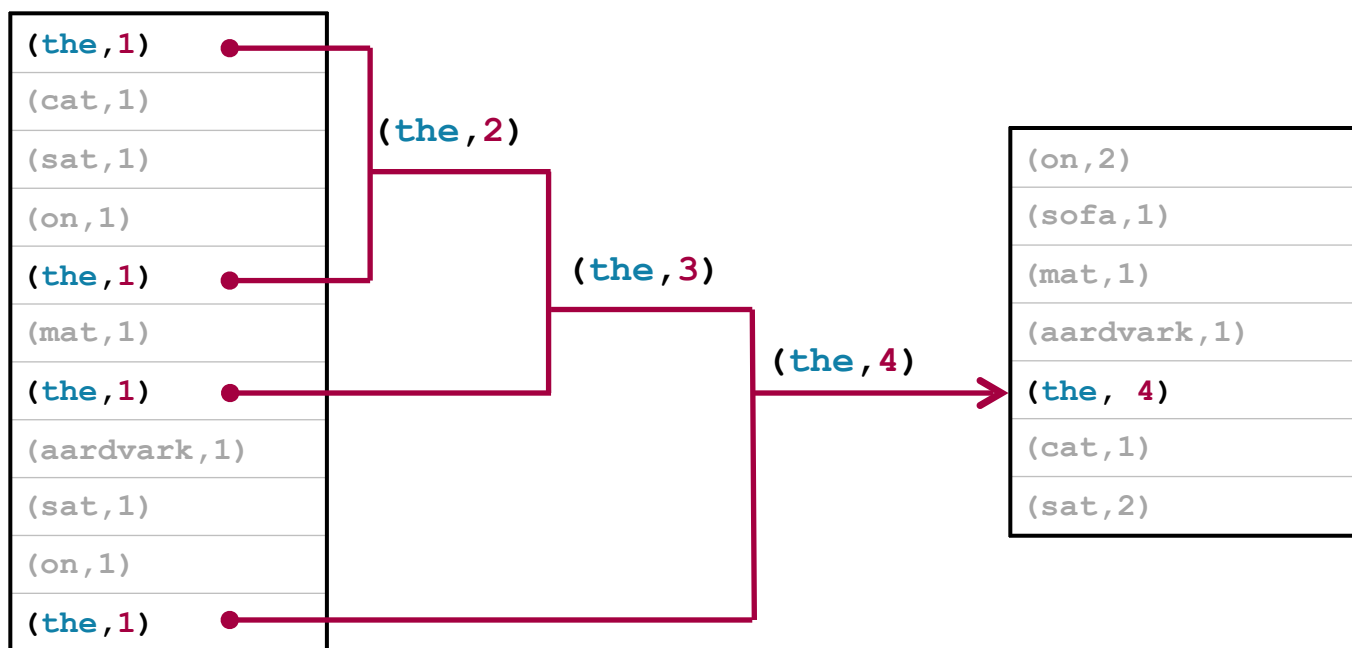
- `flatMap`, followed by `map` to map the word as key and output with a '1'
- This is the same technique you use with Hadoop MapReduce when implementing WordCount



- After `flatMap` and `map` we have key-value pairs that can be the input to a reduce
- followed by `reduceByKey` specifying addition of two values, and that sum gets added to another value, and so on:



- The function passed to `reduceByKey` combines values from two keys
 - Function must be binary




```
> val counts = sc.textFile(file) .  
  flatMap(line => line.split("\\W")) .  
  map(word => (word, 1)) .  
  reduceByKey((v1, v2) => v1 + v2)
```

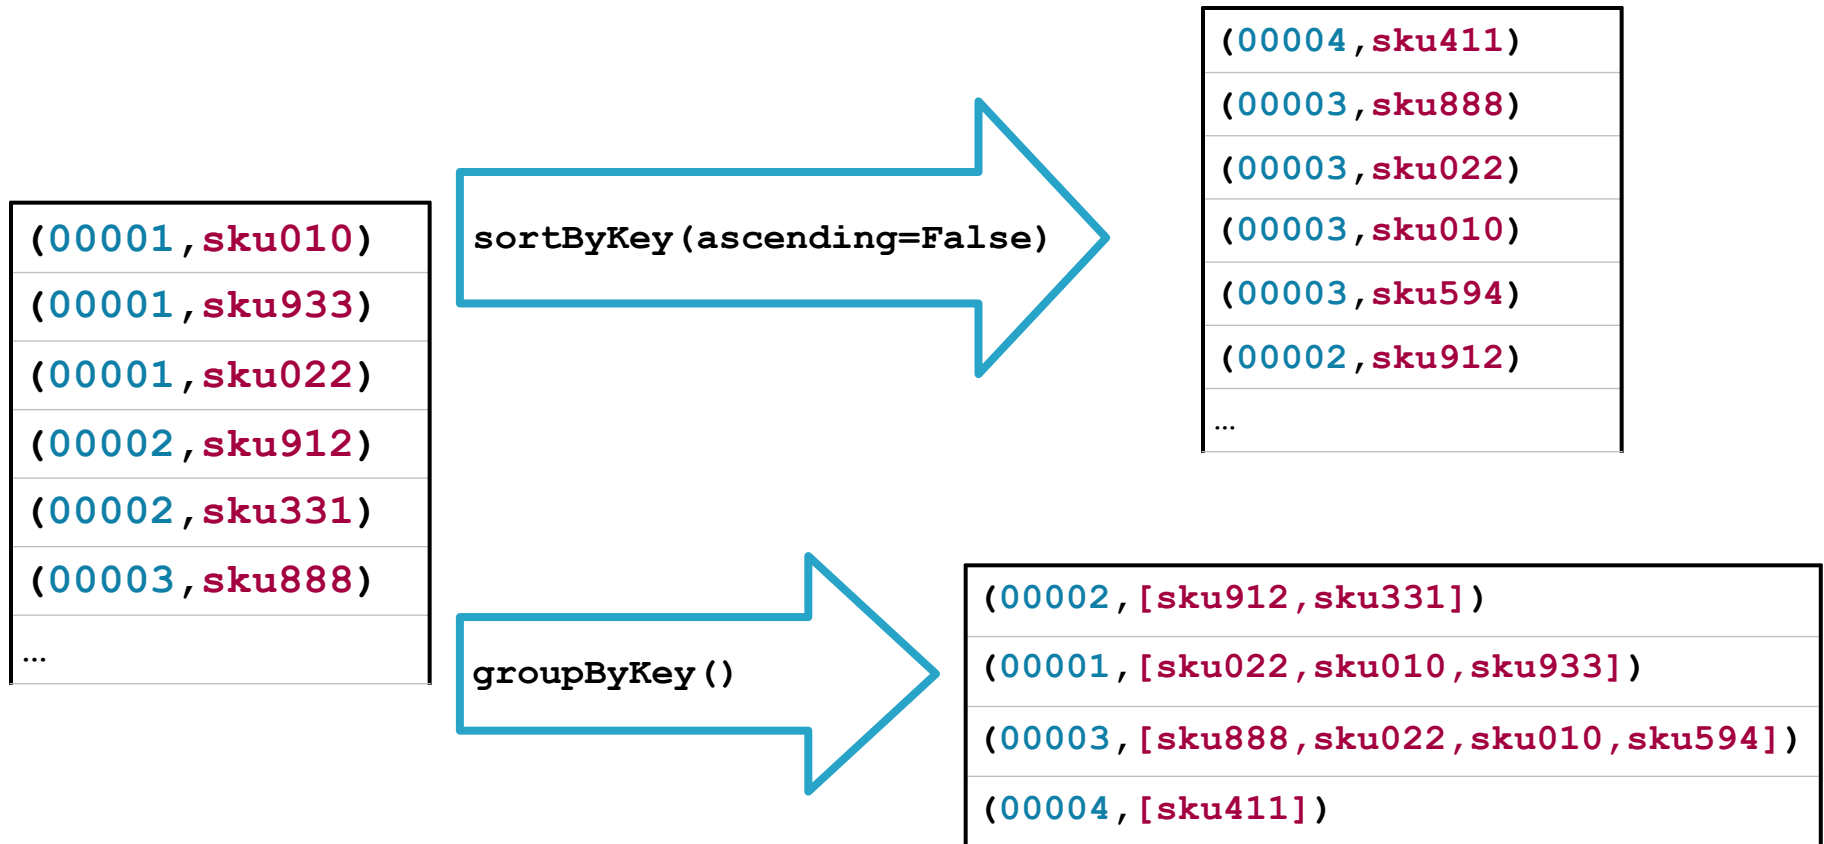
OR

```
> val counts = sc.textFile(file) .  
  flatMap(_.split("\\W")) .  
  map((_, 1)) .  
  reduceByKey(_ + _)
```

- **Word count is challenging over massive amounts of data**
 - Using a single compute node would be too time-consuming
 - Number of unique words could exceed available memory
- **Statistics are often simple aggregate functions**
 - Distributive in nature
 - e.g., max, min, sum, count
- **Map-reduce breaks complex tasks down into smaller elements which can be executed in parallel**
- **Many common tasks are very similar to word count**
 - e.g., log file analysis

- Key-Value Pair RDDs
- Map-Reduce
- **Other Pair RDD Operations**

- In addition to `map` and `reduce` functions, Spark has several operations specific to Pair RDDs
 - `countByKey`
 - Return a map with the count of occurrences of each key
 - `groupByKey`
 - Group all the values for each key in an RDD
 - `sortByKey`
 - Sort in ascending or descending order
 - `join`
 - Return an RDD containing all pairs with matching keys from two RDDs



■ Using join

RDD: movieGross

(Casablanca, \$3.7M)

(Star Wars, \$775M)

(Annie Hall, \$38M)

(Argo, \$232M)

...

RDD: movieYear

(Casablanca, 1942)

(Star Wars, 1977)

(Annie Hall, 1977)

(Argo, 2012)

...

RDD: movieGrossAndYearForEachMovie

(Casablanca, (\$3.7M, 1942))

(Star Wars, (\$775M, 1977))

(Annie Hall, (\$38M, 1977))

(Argo, (\$232M, 2012))

...

■ Some other pair RDD operations

- **keys** - return an RDD of just the keys, without the values
- **values** - return an RDD of just the values, without keys
- **lookup(*key*)** - return the value(s) for a key
- **leftOuterJoin, rightOuterJoin, fullOuterJoin** - join, including keys defined in the left, right or either RDD respectively
- **mapValues, flatMapValues** - execute a function on just the values, keeping the key the same

■ See the `PairRDDFunctions` class Scaladoc for a full list

Homework

Study for midterm exam.